



FOR HACKATHON BUILDERS

Hackathon Setup Guide

Build with 40,000+ open-source AI models. Serverless,
OpenAI-compatible, no GPUs to manage.

PROMO CODE **OPENATLAS26**


1 MONTH FREE

OPENAI-COMPATIBLE

Sign up for Feather Chat

- Scan the QR code or open the signup link, your promo code is already attached to it.
- The code OPENATLAS26 is applied automatically at checkout.
- You get one month of Feather Chat for free. If the page shows an error or the code isn't applied, reload it once.

Redeem OPENATLAS26 →



After checkout your plan shows as active, with frontier models unlocked. Cancel anytime.




SCAN HERE

STEP 02

Create your API key

- After subscribing, click your profile in the top right and open API Keys.
- Create a new key and copy it somewhere safe, you'll need it for every request.
- Treat it like a password: keep it in an environment variable, never commit it to your repo.

```
# every request carries it as a bearer token  
export FEATHERLESS_API_KEY=fw-...  
  
Authorization: Bearer $FEATHERLESS_API_KEY
```



PROFILE → API KEYS

STEP 03


Choose a model

- Browse the catalog, 40,000+ open-source models, searchable by task, size and license.
- Copy the model ID exactly as shown, it goes straight into your API call.
- Not sure where to start? The next page has five picks that cover most hackathon projects.

featherless.ai/models

model IDs look like this

deepseek-ai/DeepSeek-V3.2



FILTER BY TASK, SIZE AND LICENSE

Five models that cover most projects

DeepSeek-V3.2

Advanced reasoning and coding capabilities

MiniMax-M2.5

Excellent in agentic tool use

Kimi-K2.5

Multimodal from the ground up

Mistral-Nemo-Instruct

Fast and efficient processing

GLM-5

Excels in long horizon tasks

40,000+ more

Every size, task and license.

featherless.ai/models →

STEP 04

Your first API call

PYTHON · REQUESTS

```
import requests

response = requests.post(
    "https://api.featherless.ai/v1/chat/completions",
    headers={"Authorization": "Bearer YOUR_API_KEY"},
    json={
        "model": "deepseek-ai/DeepSeek-V3.2",
        "messages": [
            {"role": "user", "content": "Hello!"}
        ]
    }
)

print(response.json()["choices"][0]
      ["message"]["content"])
```

- Swap in the API key you created in step 02.
- The model field takes any ID from the catalog.
- The response comes back as standard JSON.
- Works the same from JavaScript, Go, curl, or any language with an HTTP client.

...or use the OpenAI SDK

PYTHON · OPENAI SDK

```
from openai import OpenAI

client = OpenAI(
    api_key="YOUR_API_KEY",
    base_url="https://api.featherless.ai/v1"
)

response = client.chat.completions.create(
    model="deepseek-ai/DeepSeek-V3.2",
    messages=[
        {"role": "user", "content": "Hello!"}
    ]
)

print(response.choices[0].message.content)
```

- Already built on OpenAI's SDK? Change `api_key` and `base_url`, keep the rest of your code.
- Streaming and tool calling work the same way.
- LangChain, LlamaIndex and other OpenAI-compatible tools plug in the same two settings.

Any OpenAI-compatible SDK works, just point `base_url` at <https://api.featherless.ai/v1>

Two endpoints to know

`/v1/chat/completions`

Best for chatbots, agents and assistants

- Structured messages with roles (system, user, assistant)
- Maintains conversation context automatically
- The endpoint you'll use for 95% of hackathon projects

`/v1/completions`

Best for raw text generation and extraction

- Takes a single prompt string
- Direct control over the prompt format
- Maximum flexibility for custom use cases

Full reference and streaming, embeddings and more: featherless.ai/docs

GOOD TO KNOW

You can run 4 requests at once

Featherless never meters tokens, send as many requests as you like. Feather Chat gives you 4 concurrent units, and each running request holds some of them until it finishes. Bigger model, more units:

MODEL SIZE	EXAMPLES	UNITS / REQUEST	ON YOUR PLAN
Small (7B – 15B)	Qwen 2.5 7B, Llama2 13B	1	4 at once
Medium (24B – 34B)	Qwen 32B Coder, Mistral 3 24B	2	2 at once
Large (70B – 72B)	Llama 3.3 70B, Qwen 2.5 72B	4	1 at once
DeepSeek-V3.2, Kimi-K2.5, GLM-5	Frontier open models	4	1 at once

Hit the limit? Extra requests get a 429, just retry as one finishes. Full details: featherless.ai/docs/concurrency-limits

Troubleshooting in 10 seconds

401 UNAUTHENTICATED

API key not recognized. Check you copied it correctly, or create a new one in account settings.

403 UNAUTHORIZED

Model is gated. Open the model's page, click "Unlock Model" and accept the license terms.

500 INTERNAL SERVER ERROR

Request could not be processed. Check for unsupported parameters in the API documentation.

503 SERVICE UNAVAILABLE

Cold model or full capacity. Retry the request; if it persists after three attempts, ping us on Discord.

THAT'S IT

Ready to build

01

Test your setup

Create your API key and run one of the code examples

02

Explore models

Find the right model for your use case in the catalog

03


Read the docs

Application guides cover common stacks and patterns

04

Start building

Ship something great, we can't wait to see it



QUICK LINKS

Redeem OPENATLAS26 →

featherless.ai/models

featherless.ai/docs

Good luck, and have fun. We can't wait to see what you build.